

# **Bitwise operators and the Memory Hierarchy**

**William Hendrix**

# Today

---

- Digital representation
- Bitwise operations
- Memory hierarchy
- Caching

# Data representation

---

- Integer data: binary numbers
  - 0, 1, 10, 11, 100, etc.
  - Unsigned `int` go up to 4B (32 bit)
  - Signed `int` represent -2B to 2B
- Decimal numbers: IEEE floating point (32 or 64 bits)
  - 1<sup>st</sup> bit: positive/negative
  - 2<sup>nd</sup>-8<sup>th</sup> bit: exponent ( $6.02 * 10^{23}$ )
  - 9<sup>th</sup>-32<sup>nd</sup> bit: mantissa (**6.02** \*  $10^{23}$ )
  - Special values to represent +/- infinity and NaN (not a number)
    - `<cmath>`: `isinf(x)`, `isnan(x)`
  - Double precision has 11-bit exponent and 52-bit mantissa
- Text data: ASCII codes
  - Codes in the range 0-255 that represent all standard characters

# Bitwise operations

---

- Operations that directly affect the bits of data
- Code is efficient but hard to read/understand
  - Use at your own risk
- Left and right shift (<< and >>)
  - Bits move left or right the given number of places
    - Precedence is very low; use () as necessary
  - **Example:**  $0x68 \ll 4 = 0x680$ ;  $0x68 \gg 1 = 0x34$
  - Since bits worth twice as much moving left, left shift effectively multiplies by  $2^n$  (barring overflow)
    - Right shift divides by  $2^n$  (integer division/floor behavior)
    - Much more efficient than “true” multiplication
  - Negative integers may copy top bit when right shifting
    - Left shift always fills in zeros
- Bitwise negation (~)
  - Flip all bits
  - **Example:**  $\sim 0xf0 = 0x0f$

# Bitwise operations (cont'd)

- Bitwise and (&) and or (|)
  - Each bit of result is the and (or) of corresponding input bits
  - **Example:**  $0x6f \ \& \ 0x30 = 0x20$ ,  $0x6f \ | \ 0x30 = 0x7f$ 

$$\begin{array}{r} 01101111 \\ \& \ 00110000 \\ \hline 00100000 \end{array}$$

$$\begin{array}{r} 01101111 \\ | \ 00110000 \\ \hline 01111111 \end{array}$$
  - Does not support short-circuit evaluation
    - Do not use in place of logical and (&&) and or (||)
- Use bitwise or to turn bits on or test
  - $x \ |= \ 0x80$  will set top bit of smallest byte
- Use bitwise and to test bits or turn bits off
  - $x \ \& \ 1$  is true (non-zero) if the last bit of  $x$  is set
  - $x \ \&= \sim 0x80$  will turn off the top bit of the last byte
- Bitwise xor (^)
  - Exclusive or (one or the other, but not both)
  - Some clever applications

# Flags

---

- Binary data can be treated as a sequence of `bool` (0/1) values
- Represent a combined state by taking bitwise or (`|`) of conditions
- Example

```
//Constants should always be powers of 2
#define TWO_DIGIT_YEAR    0x1
#define HYPHEN_SEPARATOR 0x2
#define NAME_OF_MONTH     0x4
#define PRINT_NEWLINE     0x8
void Date::print(char flags);
//Use flags & constant to test printing options

//In main():
Date today;
today.print(NAME_OF_MONTH | PRINT_NEWLINE);
```

# Negative integers

- Negative numbers represented using *two's complement*
  - Highest bit value is negated
    - **Example** (8 bits, signed):
      - `0x80`: -128
      - `0xff`: -1
  - Largest negative number: only first bit set
  - Largest positive number: all but first bit set
  - All bits set: -1
  - Flip all bits of  $x$ :  $-x - 1$
- Advantage
  - No special cases needed to add positive and negative numbers
    - Easier to design circuits to quickly add numbers
  - **Example:** `0xff + 0x01 = 0x100`
    - Carry bit is lost due to max size of data type
- **Overflow:** exceeding the max size of data
  - Occurs when carrying into/borrowing from the topmost bit
  - Eventually, integers will “wrap around” and become negative

# CPU organization

---

- All operations take place on data in registers
  - 32 or 64 registers in CPU
  - Each CPU cycle, instructions are loaded, interpreted, and run (*fetch, decode, execute*)
- Different types of instructions
  - Arithmetic
    - Addition, multiplication, division, etc. (integer and floating point)
  - Control flow
    - Decide what to do next based on comparison (equal, less than, etc.)
  - Memory operations
    - Store or retrieve information from memory/disk
- Fetch-decode-execute cycle occurs ~2B times per second
  - **Major challenge:** How do we supply data to CPU fast enough for it to keep running?



# The Memory Hierarchy

- Organization of data storage hardware
  - Higher levels are faster, more expensive, and hold less

Memory level	Access time	Capacity
Registers	~1 ns	< 1 kB
L1 cache	10-100 ns	10-50 kB
L2 cache	100-1000 ns	1-5 MB
L3 cache*	0.5-10 $\mu$ s	10-1000 MB
Main memory (RAM)	1-10 $\mu$ s	1-10 GB
Nonvolatile memory (e.g., flash)*	10-100 $\mu$ s <sup>†</sup>	100-1000 GB
Hard disk	1-10 ms	100-1000 GB
Tertiary storage (tape)*	Sec to min	1-10 TB per tape

CPU is  
here



Variables  
are here



Files are  
here


\* Not common or present on all systems

<sup>†</sup>Writing data to flash memory is much slower

- Values in table are approximate; will change over time

# Caching

- **Problem:** main memory is ~1000 times slower than CPU!
  - Files are even worse!
- **Solution:** caching
  - Store recent data in higher levels of memory hierarchy
    - E.g., *buffering* file data in memory
  - Subsequent accesses much faster
  - Retrieves a *page* of memory at once



- **Example**

```
for (int i = 0; i < nrow; i++)  
    for (int j = 0; j < ncol; j++)  
        sum += arr[i*nrow+j]; (consecutive memory access)  
for (int i = 0; i < ncol; i++)  
    for (int j = 0; j < nrow; j++)  
        sum += arr[j*nrow+i]; (random access)
```

# Hard disks

- Magnetic disks
  - Multiple *platters* per drive
  - Divided into concentric *tracks*
  - Tracks split into radial *sectors*
  - Access to outer tracks is faster
- Latency
  - Two main sources
    - Moving to correct track
  - Seek time
    - Moving to correct sector
- Throughput
  - Data after the first sector is retrieved much faster
  - Can be improved even further by reading from multiple disks at once
    - RAID

